

Mama Budget: 50-Case Scoring Protocol

A small, honest test of item-level need/want decisions and the “never scold a need” rule

Claim to test: Mama agrees with participating mothers on clear cases, while protecting every consensus need from a scolding. This is a pilot benchmark, not proof of universal judgment.

Protocol

- 1. Freeze 50 cases** — Use anonymized item, price, and only the context needed to judge it. Include clear needs, discretionary wants, context-dependent items, same-store mixed carts, and a family transfer. Do not tune the rules after seeing the labels; if you do, call these development results and test a fresh set.
- 2. Collect independent labels** — Have each available mother label each case privately as **Need**, **Want**, or **Context / unsure** before seeing Mama’s output. Three raters are ideal. A 2-of-3 label is the reference; with two raters, disagreement means “No majority.” Count every non-unanimous case.
- 3. Run the frozen rules** — Show the same item and context to Mama’s deterministic judge. Record its label, whether she actually reacted, and the rule outcome. Keep the model out of classification. Treat an unclear/no-majority item as “Ask” or quiet; do not scold it.
- 4. Freeze the comparison baseline** — Before scoring, write a simple merchant-category-to-label map using only the bank category. Apply it to the same cases. Call this the **merchant-only heuristic**, not “the bank’s judgment.” Both systems must use the same clear-case denominator.

Report these metrics separately

Measure	Calculation / interpretation
Mama agreement	Exact Need/Want matches ÷ cases with a mother majority of Need or Want. Show x/N and %, plus Need and Want rates separately.
Need-safety breaches	Consensus-Need cases where Mama reacted ÷ consensus-Need cases. Show count and rate; target is zero scolds .
Unclear handling	For Context/unsure or No-majority cases, show how often Mama chose Ask/quiet. Keep these out of the clear-case accuracy denominator.
Merchant-only match	Exact baseline matches on the same cases and denominator as Mama. Label it as a simple category-only comparison.
Rater disagreement	Number of cases without unanimous mother labels, reported out of 50. This shows where “a real mother” is not a single answer.
Stage proof	Run rice/soap (quiet), AirPods (reaction), and money sent home (quiet); then show the frozen 50-case counts. Don’t imply the demo itself proves broad accuracy.

Privacy: Use consented, anonymized examples; never send account numbers, card details, or private receipt images in the labeling chat. Nessie is the simulated bank layer for the demo, not evidence of live bank access.